

# Chapter 2: Statistical Learning

## Chapter 2: What is statistical learning

We observe a quantitative response  $Y$  and  $p$  different predictors,  $X_1, X_2, \dots, X_p$ . We assume a relationship between  $Y$  and  $X$  which can be written in the general form:

$$Y = f(X) + \epsilon \quad (2.1)$$

$f$  represents the unknown function of  $X$  and  $\epsilon$  represents the error term.  $f$  is the systematic information that  $X$  provides about  $Y$ .

### Why do we estimate $y$ ?

1. **Prediction** \* $\hat{f}$  is the estimate for  $f$  and  $\hat{Y}$  is the prediction the prediction for  $Y$ .
  - *reducible error*: inaccuracy with predictions
  - *irreducible error*:  $Y$  is a function of  $\epsilon$  and therefore cannot reduce the error introduced by  $\epsilon$

$$\hat{Y} = \hat{f}(X) \quad (2.2)$$

$$E(Y - \hat{Y})^2 = E[f(X) + \epsilon - \hat{f}(X)]^2 = [f(X) - \hat{f}(X)]^2 + Var(\epsilon)$$

The first term  $[f(X) - \hat{f}(X)]^2$  is the reducible error and the variance of epsilon is the irreducible error.

2. **Inference** \*This is where we want to estimate but not predict  $f$ .
  - Which predictors are associated with the response?
  - What is the relationship between the response and each predictor?
  - Can the relationship between  $Y$  and each predictor be adequately summarized using a linear equation?

### How do we estimate $f$ ?

*Training data* is a set of observations used to teach a method to estimate  $f$ .

We can use two models to estimate  $f$ .

**Parametric models** These models have a two step approach: - First Assume  $f$ :

$$f(X) = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p \quad (2.4)$$

This is a linear model where instead of having to estimate an entirely arbitrary  $p$ -dimensional function  $f(X)$  we only need to estimate  $p + 1$  coefficients. \* fit/train the model

$$Y \approx \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_p X_p$$

We need to estimate the parameter such that the approximate  $Y$ . This is called *ordinary least squares*.

This procedure is called parametric because estimating  $f$  is through the estimation of a set of parameters. While much simpler to do, this can cause approximation variance.

*Flexible models* can fit many functional forms for  $f$ . *Over fitting* can occur if the model is too complex and follows *noise* too closely.

**Nonparametric models** Nonparametric models don't make explicit assumptions but seek estimates of  $f$ . While we don't need to worry about fit of data, a large amount of data is needed to obtain an accurate estimate of  $f$ .

*Thin-plate spline* is used to estimate  $f$  by not imposing a predefined model on  $f$  but attempts to produce estimates for  $f$  that is as close to observed data.

This observed fit is called *smoothness*. The smoother the visualization of the data is, the better fit. Choosing the correct metric of smoothing to avoid over-fitting will be discussed.

### Extra: Smoothness measure

*Data fidelity term* The TPS smoothness measure arises from considering the integral of the second derivative (usually denoting the curvature/concavity of a graph). In this case, we can use the *energy function* (a function we want to minimize/maximize):

$$E_{tps}(f) = \sum_{i=1}^K ||y_i - f(x_i)||^2$$

Known as the *data fidelity term*, this measures how well function  $f$  fits the data points  $(x_i, y_i)$ . The goal is to minimize the difference or error of the observed and true values of  $y_i$  and  $f(x_i)$ .

*Smoothness Penalty Term* The smoothness variant uses a tuning parameter  $\lambda$  (which is externally constructed) to control the "rigidity" and minimize  $E_{tps,smooth}$  with a unique minimizer  $f$ :

$$E_{tps,smooth}(f) = \sum_{i=1}^K ||y_i - f(x_i)||^2 + \lambda \int \int [(\frac{\partial^2}{\partial x_1^2})^2 + 2(\frac{\partial f}{\partial x_1 x_2})^2 + (\frac{\partial^2 f}{\partial x_2^2})^2] dx_1 dx_2$$

This penalty term includes the sum of squared second order partial derivatives that quantify the curvature of  $f$  in different directions such as: \*  $\frac{\partial^2}{\partial x_1^2}$ : curvature in  $x_1$  direction \*  $\frac{\partial f}{\partial x_1 x_2}$  mixed curvature between  $x_1$  and  $x_2$  \*  $\frac{\partial^2 f}{\partial x_2^2}$  curvature in  $x_2$  direction

### The trade-off between prediction accuracy and model interpretability

Thin plate splines can generate much wider ranges of possible shapes to estimate  $f$ .

A more restrictive model is used for inference whereas linear models are good for prediction. Splines can lead to super complicated estimates of  $f$  that lead to difficult predictor separation to analyze the effect on the response.

**Other models** Restrictive - linear regression - thin plate splines - lasso - GAM Flexible - thin plate splines - bootstrapping - bagging, boosting, support vector machines

### Supervised vs unsupervised learning

*Supervised learning* is used when we wish to fit a model that relates the response to the predictors with the aim of accurately predicting the response for future observations. Examples include - linear/logistic regression - GAM - boosting, support vector machines




Figure 1: Representation of tradeoff

*Unsupervised learning* is used when we observe a vector of measurements  $x_i$  but no associated response  $y_i$ . We lack a response to supervise our analysis so instead we understand relationships between variables and observations. Examples - cluster analysis

*Semi supervised learning* is used when we have disjoint sets of supervised and unsupervised observations. We wish to incorporate  $m$  observations which response measurements are available and  $n - m$  observations for which they are not.

## Regression vs. classification problems

Variables can be characterized as either *quantitative* or *qualitative/categorical*.

We tend to prefer regression for quantitative (ex: least squares regression) and classification for categorical variables (ex: logistic regression). Some methods like K-nearest neighbors and boosting used for both cases.

## 2.2 Assessing model accuracy

### Measuring the quality of fit

We need to quantify the extent to which the predicted response value for a given observation is close to the true response value for that observation.

This is called the *Mean Squared Error*

$$MSE = \frac{1}{n} \sum_{i=1}^n (y_i - \hat{f}(x_i))^2 \quad (2.5)$$

$\hat{f}(x_i)$  is the prediction that  $\hat{f}$  gives for the  $i$ th observation. The MSE is computed from the training data.

*training MSE* is used when we are interested in the accuracy of the predictions that we obtain when we apply our method to previously unseen test data.

Described mathematically: suppose we use our training observations  $\{(x_1, y_1), \dots, (x_n, y_n)\}$  to estimate  $\hat{f}$ . We can then compute  $\hat{f}(x_1), \dots, \hat{f}(x_n)$ . If our response is close to the true response, then our MSE is small. But we don't want to know if  $\hat{f}(x_i) \approx y_i$ .


We want to know whether  $\hat{f}_0$  is approx equal to  $y_0$  where  $(x_0, y_0)$  is previously unseen data. We need to choose the method that gives the lowest *test MSE* as opposed to the lowest training MSE.

We can use the *average squared prediction error*

$$Ave(y_0 - \hat{f}(x_0))^2 \quad (2.6)$$

We want the method with the smallest average test MSE.

What if we don't have test observations to find the smallest test MSE? We can't use the training MSE because there is no guarantee that the method with the lowest training MSE will also have the lowest test MSE. The problem is that many methods estimate coefficients to minimize the training set and ignore testing measures.



The true  $f$  is approx linear. We observe the training MSE decreases monotonically as the model flexibility increases. But because the truth is closer to linear, the test MSE only slightly decreases before increasing again. Therefore linear regression is better than splicing.

The flexibility level corresponding to the model with the minimal test MSE can vary considerably among data sets. Using *cross validation* can be used to estimate the test MSE using the training data.

### Bias-Varinace trade off

The U-shapes MSE's is the consequence of two competing properties of statistical learning methods. It can be shown that the expected test MSE for a given value  $x_0$  can always be decomposed into the sum of three fundamental quantities: - variance of  $\hat{f}(x_0)$  - squared bias of  $\hat{f}(x_0)$  - variance of error term  $\epsilon$

$$E(y_0 - \hat{f}(x_0))^2 = Var(\hat{f}(x_0)) + [Bias(\hat{f}(x_0))]^2 + Var(\epsilon) \quad (2.7)$$

$E(y_0 - \hat{f}(x_0))^2$  is the *expected test MSE* and refers to the average test MSE that we would obtain if we repeatedly estimated  $f$  using large number of training sets. - the overall expected MSE can be computed by averaging the expected test MSE over all possible values of  $x_0$  in the test set

2.7 shows that to minimize the expected test MSE we need to minimize variance and bias.  $Var(\epsilon)$  is non-zero and irreducible, therefore the expected test MSE can never be smaller than the variance of  $\epsilon$ .

*Variance* is the amount in which  $\hat{f}$  would change if using a different training data set. *Bias* is the error that is introduced by approx real life problem. Generally more flexible methods result in less bias as the variance increases while the bias decreases. The rate of change between the two determines the increase or decrease of the test MSE.

### The classification setting

The common approach of quantifying the accuracy of our estimate with categorical data is the training *error rate*, which is the proportion of mistakes that are made if we apply our estimate  $\hat{f}$  to the training observations:

$$\frac{1}{n} \sum_{i=1}^n I(y_i \neq \hat{y}_i) \quad (2.8)$$

$I(y_i \neq \hat{y}_i)$  is the indicator variable that equals 1 if  $y_i \neq \hat{y}_i$  and 0 otherwise. 0 means that the variable was classified correctly and vice versa.

2.8 is the training error because it is computed based on the training data. We are most interested in the average *test* error over the set of *test observations* of the form  $(x_0, y_0)$  given by

$$Ave(y_i \neq \hat{y}_i) \quad (2.9)$$

### The Bayes Classifier It is possible to show that the test error rate given in 2.9 is minimized on average by a simple classifier that assigns each observation to the most likely case given its predictor values.

In other words: simply assign a test observation with predictor value  $x_0$  to the class  $j$  for which

$$Pr(Y = j | X = x_0) \quad (2.10)$$

*The proof behind the pudding* Citation: Shuzhan Fan **The Bayes theorem** Given a feature vector  $X = (x_1, \dots, x_n)$  and a class variable  $C_k$ , the Bayes theorem states that:

$$P(C_k | X) = \frac{P(X | C_k) P(C_k)}{P(X)} \text{ for } k = 1, 2, \dots, K$$

$P(C_k | X)$  is the posterior probability and  $P(X | C_k)$  is the likelihood.  $P(C_k)$  is the prior probability of class and  $P(X)$  is the probability of the predictor.

Using the chain rule the likelihood  $P(X | C_k)$  can be decomposed as:

$$P(X | C_k) = P(x_1 \dots x_n | C_k) = P(x_1 | x_2, \dots, x_n, C_k) P(x_2 | x_3, \dots, x_n, C_k) \dots P(x_{n-1} | x_n, C_k) P(x_n | C_k)$$

### Naive independence assumption

The naive conditional independence assumption allows us to calculate  $P(X | C_k)$  easily by:

$$P(x_i | x_{i+1}, \dots, x_n | C_k) = P(x_i | C_k)$$

We can get:

$$P(X | C_k) = P(x_1, \dots, x_n | C_k) = \prod_{i=1}^n P(x_i | C_k)$$

The posterior probability  $P(C_k | X)$  can be written now as:

$$P(C_k | X) = \frac{P(C_k) \prod_{i=1}^n P(x_i | C_k)}{P(X)}$$

### Naive Bayes model Because the predictor  $P(X)$  is constant given the input, we can get  $P(C_k | X)$  which is positively proportional to:

$$P(C_k | X) \propto P(C_k) \prod_{i=1}^n P(x_i | C_k)$$

Therefore the naive Bayes classification for different class values of  $C_k$  looks at the maximum of  $P(C_k) \prod_{i=1}^n P(x_i | C_k)$  which can be formulated to:

$$\hat{C} = \operatorname{argmax}_{C_k} P(C_k) \prod_{i=1}^n P(x_i | C_k)$$


The prior probability of class  $P(C_k)$  could be calculated as the relative frequency of class  $C_k$  in the training data.

**The difference between naive Bayes vs Bayes** The likelihood  $P(x_i|C_k)$  is usually modeled using the same class distributions, but for the naive Bayes classifier the assumption is made regarding the distribution of  $P(x_i|C_k)$ .

A conditional probability is represented as the probability of  $Y = j$  given the observed predictor vector  $x_0$ . This complete classifier is called the *Bayes classifier*.

In a two class problem where there is only class 1 and class 2, the Bayes classifier categorizes classes by a parameter like  $P(Y = 1|X = x_0) > 0.5$ .

The *Bayes decision boundary* is the line that represents the points where the probability is exactly 50%. The Bayes classifier prediction is determined by the Bayes decision boundary. The image below illustrates this:



The Bayes classifier produces the lowest possible test error rate, called the *Bayes error rate*. Because the Bayes classifier will always choose the class which the conditional probability is the largest, the error rate  $X = x_0$  will be  $1 - \max_j P(Y = j|X = x_0)$ .

The overall Bayes error is given by

$$1 - E(\max_j P(Y = j|X)) \quad (2.11)$$

Because the classes overlap in the true population in the image above,  $\max_j P(Y = j|X = x_0) < 1$  for some values of  $x_0$ . The *Bayes error rate is analogous to the irreducible error*.

## K-nearest neighbors

For real data we don't have conditional distributions of  $Y$  given  $X$ . so computing the Bayes classifier is impossible. Many approaches attempt to estimate Bayes estimator as the gold standard and then classify a given observation to the class with the highest *estimated* probability.

A method that does this is called the KNN classifier. Given a positive integer  $K$  and a test observation  $x_0$ , the KNN classifier first identifies the  $K$  points in the training data that are closest to  $x_0$  represented by  $\mathcal{N}$ .

It estimates the conditional probability for the class  $j$  as the fraction of points in  $\mathcal{N}_0$  whose response values equal  $j$ :

$$P(Y = j|X = x_0) = \frac{1}{K} \sum_{i \in \mathcal{N}_0} I(y_i = j) \quad (2.12)$$

Finally the KNN applies Bayes rule and classifies the test observation  $x_0$  to the class with the largest probability.

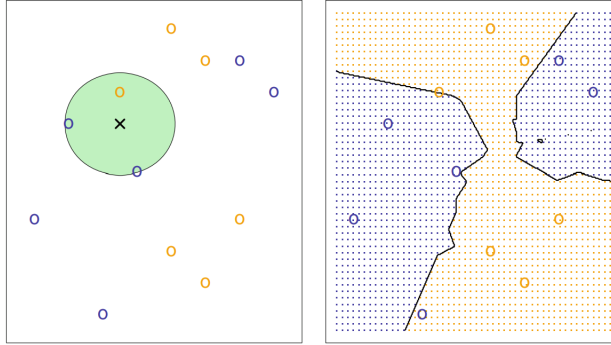


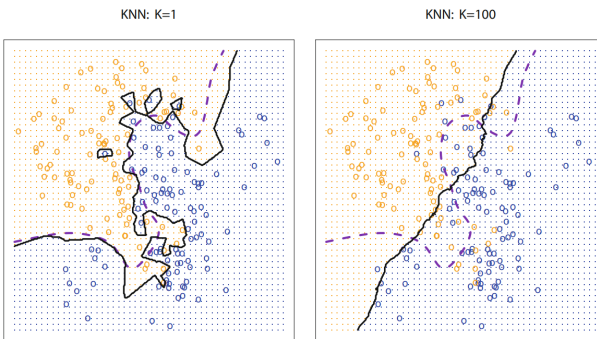
Figure 2: The KNN Classifier

In the example below we set  $K = 3$  at all possible values for  $X_1$  and  $X_2$ . The KNN decision boundary is illustrated in green. The black cross illustrates the test observation at which a predicted class label is desired. - the three closest points are identified and classified to blue: the most commonly occurring class


The KNN boundary in black splits blue and orange class observation regions.

**The choice of  $K$**  The choice of  $K$  has a drastic effect on the KNN classifier. When  $K = 1$  the decision boundary is overly flexible and over fits the data. Logically, this corresponds to a classifier that has low bias but very high variance. As  $K$  grows the method becomes less flexible and produces a decision boundary that is close to *linear*. - *this gives us a low variance but high bias classifier*

**In general as we use more flexible classification methods the training error rate will decline but the test error rate may not.**



With  $K = 1$  the decision boundary is too flexible and vice versa for  $K = 100$ .



The black dashed line is the Bayes error rate. The error rates are measured as the level of flexibility, accessed by  $\frac{1}{K}$ , increases, (ie the number of neighboring  $K$ s decrease). *The more flexible the less neighboring  $K$ s*

## Excercises

1.

- a) flexible
- b) inflexible
- c) inflexible
- d) flexible

Usually, inflexible methods are used for predicting with large numbers of predictor  $p$  or with a strong relationship. Flexible methods are for large sample sizes that correspond to high variance but low bias. In a case with a large  $n$  and small  $p$ , we can use flexible methods without the fear of overfitting.

2.

- a) regression, inference
- b) classification, inference
- c) regression, prediction

Applied

```
# read csv
college <- read.csv("C:/Users/chuan_71/OneDrive/Desktop/TRADING/ISL/College.csv")

# data exploration
head(college)
```

```
##              X Private Apps Accept Enroll Top10perc Top25perc
## 1 Abilene Christian University    Yes 1660   1232    721      23      52
## 2              Adelphi University    Yes 2186   1924    512      16      29
## 3                Adrian College    Yes 1428   1097    336      22      50
## 4            Agnes Scott College    Yes  417    349    137      60      89
```



```
## 5      Alaska Pacific University      Yes 193    146    55      16      44
## 6      Albertson College              Yes 587    479   158      38      62
##      F.Undergrad P.Undergrad Outstate Room.Board Books Personal PhD Terminal
## 1      2885          537      7440          3300    450      2200  70      78
## 2      2683          1227     12280          6450    750      1500  29      30
## 3      1036           99      11250          3750    400      1165  53      66
## 4       510           63      12960          5450    450       875  92      97
## 5       249          869      7560          4120    800      1500  76      72
## 6       678          41      13500          3335    500       675  67      73
##      S.F.Ratio perc.alumni Expend Grad.Rate
## 1      18.1          12      7041          60
## 2      12.2          16     10527          56
## 3      12.9          30      8735          54
## 4       7.7          37     19016          59
## 5      11.9           2     10922          15
## 6       9.4          11      9727          55
```

```
college = college[,-1]
```

```
summary(college)
```

```
##      Private              Apps              Accept              Enroll
## Length:777      Min.   : 81      Min.   : 72      Min.   : 35
## Class :character 1st Qu.: 776      1st Qu.: 604      1st Qu.: 242
## Mode  :character Median : 1558      Median : 1110      Median : 434
##              Mean   : 3002      Mean   : 2019      Mean   : 780
##              3rd Qu.: 3624      3rd Qu.: 2424      3rd Qu.: 902
##              Max.   :48094      Max.   :26330      Max.   :6392
##      Top10perc      Top25perc      F.Undergrad      P.Undergrad
## Min.   : 1.00      Min.   : 9.0      Min.   : 139      Min.   : 1.0
## 1st Qu.:15.00      1st Qu.: 41.0      1st Qu.: 992      1st Qu.: 95.0
## Median :23.00      Median : 54.0      Median : 1707      Median : 353.0
## Mean   :27.56      Mean   : 55.8      Mean   : 3700      Mean   : 855.3
## 3rd Qu.:35.00      3rd Qu.: 69.0      3rd Qu.: 4005      3rd Qu.: 967.0
## Max.   :96.00      Max.   :100.0      Max.   :31643      Max.   :21836.0
##      Outstate      Room.Board      Books              Personal
## Min.   : 2340      Min.   :1780      Min.   : 96.0      Min.   : 250
## 1st Qu.: 7320      1st Qu.:3597      1st Qu.: 470.0      1st Qu.: 850
## Median : 9990      Median :4200      Median : 500.0      Median :1200
## Mean   :10441      Mean   :4358      Mean   : 549.4      Mean   :1341
## 3rd Qu.:12925      3rd Qu.:5050      3rd Qu.: 600.0      3rd Qu.:1700
## Max.   :21700      Max.   :8124      Max.   :2340.0      Max.   :6800
##      PhD              Terminal      S.F.Ratio      perc.alumni
## Min.   : 8.00      Min.   : 24.0      Min.   : 2.50      Min.   : 0.00
## 1st Qu.: 62.00      1st Qu.: 71.0      1st Qu.:11.50      1st Qu.:13.00
## Median : 75.00      Median : 82.0      Median :13.60      Median :21.00
## Mean   : 72.66      Mean   : 79.7      Mean   :14.09      Mean   :22.74
## 3rd Qu.: 85.00      3rd Qu.: 92.0      3rd Qu.:16.50      3rd Qu.:31.00
## Max.   :103.00      Max.   :100.0      Max.   :39.80      Max.   :64.00
##      Expend      Grad.Rate
## Min.   : 3186      Min.   : 10.00
## 1st Qu.: 6751      1st Qu.: 53.00
## Median : 8377      Median : 65.00
## Mean   : 9660      Mean   : 65.46
```

```
## 3rd Qu.:10830 3rd Qu.: 78.00
## Max. :56233 Max. :118.00
```

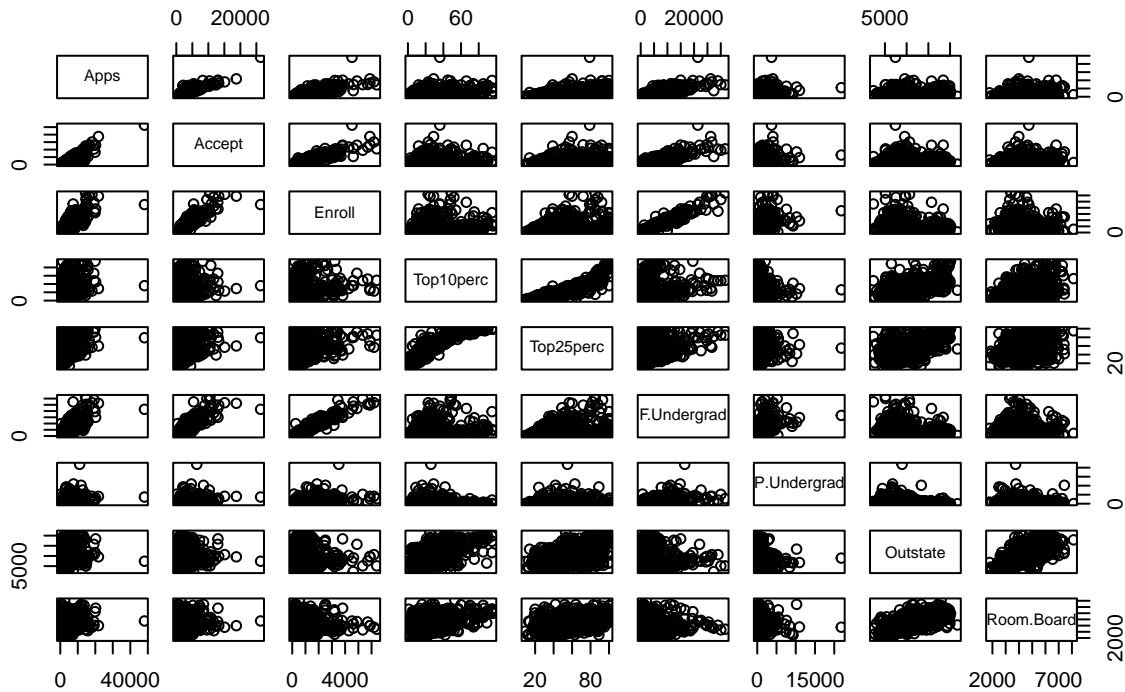
```
head(college)
```

```
## Private Apps Accept Enroll Top10perc Top25perc F.Undergrad P.Undergrad
## 1 Yes 1660 1232 721 23 52 2885 537
## 2 Yes 2186 1924 512 16 29 2683 1227
## 3 Yes 1428 1097 336 22 50 1036 99
## 4 Yes 417 349 137 60 89 510 63
## 5 Yes 193 146 55 16 44 249 869
## 6 Yes 587 479 158 38 62 678 41
## Outstate Room.Board Books Personal PhD Terminal S.F.Ratio perc.alumni Expend
## 1 7440 3300 450 2200 70 78 18.1 12 7041
## 2 12280 6450 750 1500 29 30 12.2 16 10527
## 3 11250 3750 400 1165 53 66 12.9 30 8735
## 4 12960 5450 450 875 92 97 7.7 37 19016
## 5 7560 4120 800 1500 76 72 11.9 2 10922
## 6 13500 3335 500 675 67 73 9.4 11 9727
## Grad.Rate
## 1 60
## 2 56
## 3 54
## 4 59
## 5 15
## 6 55
```

```
# matrix scatterplot
par(mfrow = c(4, 5), mar = c(4, 4, 2, 1))

pairs(college[, 2:10], main = "Scatterplot")
```

## Scatterplot



```
# side by side plot
par(mfrow = c(4, 5), mar = c(4, 4, 2, 1))

boxplot(Outstate ~ Private, data = college, main = "Boxplots of Outstate v Private", xlab = "Private", ylab = "Outstate")

# Add Elite
Elite = rep("No", nrow(college))
Elite[college$Top10perc>50]="Yes"
Elite = as.factor(Elite)
college = data.frame(college, Elite)
summary(college)
```


```
##      Private      Apps      Accept      Enroll
## Length:777      Min.   :  81      Min.   :  72      Min.   :  35
## Class :character 1st Qu.: 776      1st Qu.: 604      1st Qu.: 242
## Mode  :character Median : 1558      Median : 1110      Median : 434
##                      Mean  : 3002      Mean  : 2019      Mean  : 780
##                      3rd Qu.: 3624      3rd Qu.: 2424      3rd Qu.: 902
##                      Max.   :48094      Max.   :26330      Max.   :6392
##      Top10perc      Top25perc      F.Undergrad      P.Undergrad
## Min.   : 1.00      Min.   :  9.0      Min.   : 139      Min.   :  1.0
## 1st Qu.:15.00      1st Qu.: 41.0      1st Qu.: 992      1st Qu.: 95.0
## Median :23.00      Median : 54.0      Median : 1707      Median : 353.0
## Mean   :27.56      Mean   : 55.8      Mean   : 3700      Mean   : 855.3
## 3rd Qu.:35.00      3rd Qu.: 69.0      3rd Qu.: 4005      3rd Qu.: 967.0
```

```
## Max. :96.00 Max. :100.0 Max. :31643 Max. :21836.0
## Outstate Room.Board Books Personal
## Min. : 2340 Min. :1780 Min. : 96.0 Min. : 250
## 1st Qu.: 7320 1st Qu.:3597 1st Qu.: 470.0 1st Qu.: 850
## Median : 9990 Median :4200 Median : 500.0 Median :1200
## Mean :10441 Mean :4358 Mean : 549.4 Mean :1341
## 3rd Qu.:12925 3rd Qu.:5050 3rd Qu.: 600.0 3rd Qu.:1700
## Max. :21700 Max. :8124 Max. :2340.0 Max. :6800
## PhD Terminal S.F.Ratio perc.alumni
## Min. : 8.00 Min. : 24.0 Min. : 2.50 Min. : 0.00
## 1st Qu.: 62.00 1st Qu.: 71.0 1st Qu.:11.50 1st Qu.:13.00
## Median : 75.00 Median : 82.0 Median :13.60 Median :21.00
## Mean : 72.66 Mean : 79.7 Mean :14.09 Mean :22.74
## 3rd Qu.: 85.00 3rd Qu.: 92.0 3rd Qu.:16.50 3rd Qu.:31.00
## Max. :103.00 Max. :100.0 Max. :39.80 Max. :64.00
## Expend Grad.Rate Elite
## Min. : 3186 Min. : 10.00 No :699
## 1st Qu.: 6751 1st Qu.: 53.00 Yes: 78
## Median : 8377 Median : 65.00
## Mean : 9660 Mean : 65.46
## 3rd Qu.:10830 3rd Qu.: 78.00
## Max. :56233 Max. :118.00
```

```
# side by side plots
```


```
par(mfrow = c(4, 5), mar = c(4, 4, 2, 1))
```

```
>xplots of Outstate \
```



```
boxplot(Outstate ~ Elite, data = college, main = "Boxplots of Outstate v Elite", xlab = "Elite", ylab =
# histograms
par(mfrow = c(4, 5), mar = c(4, 4, 2, 1))
```

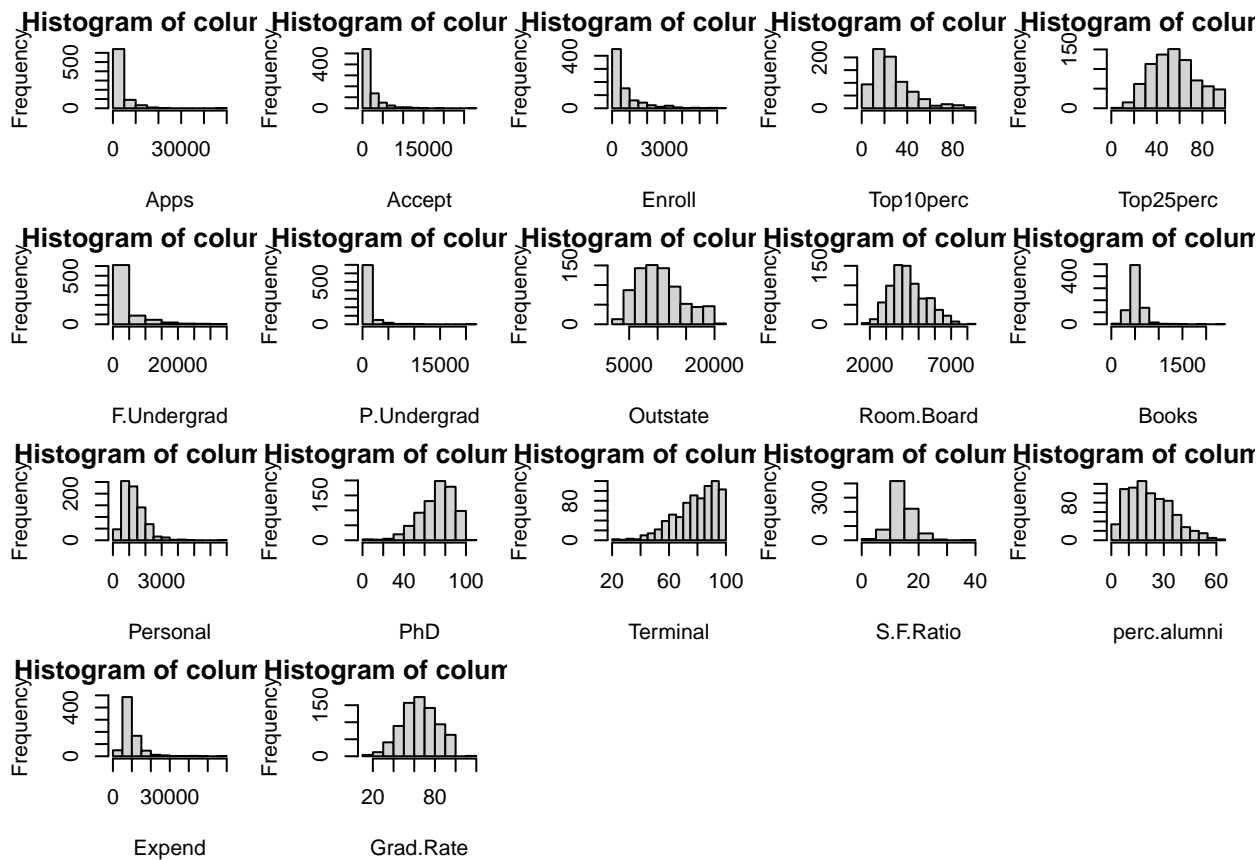
### Boxplots of Outstate



```
numeric_cols <- which(sapply(college, is.numeric))
print(numeric_cols)
```

```
##      Apps      Accept      Enroll      Top10perc      Top25perc F.Undergrad
##         2         3         4         5         6         7
## P.Undergrad      Outstate      Room.Board      Books      Personal      PhD
##         8         9        10        11        12        13
##      Terminal      S.F.Ratio      perc.alumni      Expend      Grad.Rate
##        14        15        16        17        18
```

```
for (col in numeric_cols) {
  hist(college[, col], main = paste("Histogram of column ", col), xlab = names(college)[col])
}
```



9.

```
auto <- read.csv("C:/Users/chuan_71/OneDrive/Desktop/TRADING/ISL/Auto.csv")
```

```
head(auto)
```

```
##   mpg cylinders displacement horsepower weight acceleration year origin
## 1  18         8          307         130   3504          12.0    70     1
## 2  15         8          350         165   3693          11.5    70     1
## 3  18         8          318         150   3436          11.0    70     1
## 4  16         8          304         150   3433          12.0    70     1
## 5  17         8          302         140   3449          10.5    70     1
## 6  15         8          429         198   4341          10.0    70     1
##
##           name
## 1 chevrolet chevelle malibu
## 2      buick skylark 320
## 3    plymouth satellite
## 4      amc rebel sst
## 5        ford torino
## 6      ford galaxie 500
```

```
# remove nan
auto <- na.omit(auto)
head(auto)
```

```
##   mpg cylinders displacement horsepower weight acceleration year origin
## 1  18         8          307         130   3504          12.0   70     1
## 2  15         8          350         165   3693          11.5   70     1
## 3  18         8          318         150   3436          11.0   70     1
## 4  16         8          304         150   3433          12.0   70     1
## 5  17         8          302         140   3449          10.5   70     1
## 6  15         8          429         198   4341          10.0   70     1
##                                     name
## 1 chevrolet chevelle malibu
## 2      buick skylark 320
## 3    plymouth satellite
## 4      amc rebel sst
## 5      ford torino
## 6      ford galaxie 500
```

```
# quant vs qual predictors
quant_cols <- which(sapply(auto, is.numeric))
qual_cols  <- which(!sapply(auto, is.numeric))
print(quant_cols)
```

```
##      mpg      cylinders displacement      weight acceleration      year
##      1              2              3              5              6              7
##      origin
##      8
```

```
print(qual_cols)
```

```
## horsepower      name
##          4          9
```

```
# function applicator function

app <- function(col_name, func){
  # init matrix
  results <- matrix(nrow = length(col_name), ncol = 3)
  rownames(results) <- col_name
  colnames(results) <- c("Mean", "Sd", "Range")

  # Loop
  for(i in seq_along(col_name)){
    col <- col_name[i]
    col_data <- auto[[col]]

    x <- mean(col_data)
    y <- sd(col_data)
    z <- paste(range(col_data), collapse = ", ")

    results[i, "Mean"] <- x
    results[i, "Sd"] <- y
    results[i, "Range"] <- z
  }
}
```

```

# return result
return(results)
}

```

```

# mean, sd, range
app(quant_cols)

```

```

##      Mean                Sd              Range
## 1 "23.5158690176322" "7.82580392894656" "9, 46.6"
## 2 "5.45843828715365" "1.70157698079185" "3, 8"
## 3 "193.53274559194" "104.37958329993" "68, 455"
## 5 "2970.26196473552" "847.904119489725" "1613, 5140"
## 6 "15.5556675062972" "2.74999529297615" "8, 24.8"
## 7 "75.9949622166247" "3.69000490146168" "70, 82"
## 8 "1.57430730478589" "0.802549495797039" "1, 3"

```

```

# remove 10th and 85th
auto <- auto[-c(10, 85), ]
new_quant_cols <- which(sapply(auto, is.numeric))
new_qual_cols <- which(!sapply(auto, is.numeric))

```

```

# range, mean, std
app(new_quant_cols)

```

```

##      Mean                Sd              Range
## 1 "23.5286075949367" "7.83192521828543" "9, 46.6"
## 2 "5.45569620253165" "1.69948834009909" "3, 8"
## 3 "193.279746835443" "104.061131705913" "68, 455"
## 5 "2970.23797468354" "847.764300207639" "1613, 5140"
## 6 "15.5711392405063" "2.73349731700723" "8, 24.8"
## 7 "76.020253164557" "3.68142475960302" "70, 82"
## 8 "1.57215189873418" "0.800846111706057" "1, 3"

```

10.